

This document explains the logic behind the hyperparameter tuning process performed on the `DecisionTreeClassifier`, `RandomForestClassifier`, and `AdaBoostClassifier` models, how the combinations of hyperparameters were tested, and provides an analysis of the results to determine the best-performing model.

1. Hyperparameter Tuning Logic

1.1 Decision Tree Classifier

The hyperparameter grid for the `DecisionTreeClassifier` included:

- `max_depth`: Limits the maximum depth of the tree. Values tested: `[3, 5, 10, None]`.
- `min_samples_split`: Minimum number of samples required to split an internal node. Values tested: `[2, 5, 10]`.
- `criterion`: The function to measure the quality of a split. Values tested: `['gini', 'entropy']`.

Total Combinations:

$$4 (\text{max_depth}) \times 3 (\text{min_samples_split}) \times 2 (\text{criterion}) = 24 \text{ combinations}$$

Using 5-fold cross-validation, each combination was tested 5 times, resulting in:

$$24 \text{ combinations} \times 5 \text{ folds} = 120 \text{ fits.}$$

The grid search identified the best hyperparameters:

- `criterion='entropy'`
- `max_depth=3`
- `min_samples_split=2`

1.2 Random Forest Classifier

The hyperparameter grid for the `RandomForestClassifier` included:

- `n_estimators`: Number of trees in the forest. Values tested: `[50, 100, 200]`.
- `max_depth`: Maximum depth of each tree. Values tested: `[3, 5, 10, None]`.
- `max_features`: Number of features to consider when looking for the best split. Values tested: `['sqrt', 'log2']`.
- `min_samples_split`: Minimum samples required to split an internal node. Values tested: `[2, 5, 10]`.
- `class_weight`: Strategy to handle imbalanced data. Values tested: `[None, 'balanced']`.

Total Combinations:

$$3 (\text{n_estimators}) \times 4 (\text{max_depth}) \times 2 (\text{max_features}) \\ \times 3 (\text{min_samples_split}) \times 2 (\text{class_weight}) = 144 \text{ combinations}$$

Using 5-fold cross-validation, each combination was tested 5 times, resulting in:

$$144 \text{ combinations} \times 5 \text{ folds} = 720 \text{ fits}$$

The grid search identified the best hyperparameters:

- `n_estimators=50`
 - `max_depth=10`
 - `max_features='log2'`
 - `min_samples_split=2`
 - `class_weight='balanced'`
-

1.3 AdaBoost Classifier

The hyperparameter grid for the `AdaBoostClassifier` included:

- `n_estimators`: Number of weak learners (iterations). Values tested: `[50, 100, 200]`.
- `learning_rate`: Shrinks the contribution of each weak learner. Values tested: `[0.01, 0.1, 1]`.
- `estimator`: The base learner used for boosting. Decision trees with depths `[1, 2, 3]` were tested.

Total Combinations:

$$3 (n_estimators) \times 3 (learning_rate) \times 3 (estimator \ depths) = 27 \text{ combinations}$$

Using 5-fold cross-validation, each combination was tested 5 times, resulting in:

$$27 \text{ combinations} \times 5 \text{ folds} = 135 \text{ fits.}$$

The grid search identified the best hyperparameters:

- `n_estimators=200`
 - `learning_rate=1`
 - `estimator=DecisionTreeClassifier(max_depth=3)`
-

2. Model Evaluation

Each model was evaluated on the test set using the following metrics:

- **Accuracy**: Percentage of correctly classified samples.
 - **Precision**: Proportion of true positives among all positive predictions.
 - **Recall**: Proportion of true positives among all actual positives.
 - **F1 Score**: Harmonic mean of precision and recall.
 - **Confusion Matrix**: Breakdown of true positives, true negatives, false positives, and false negatives.
-

2.1 Decision Tree Classifier

- **Confusion Matrix:** $\begin{bmatrix} 38 & 4 \\ 2 & 70 \end{bmatrix}$
 - **Metrics:**
 - Accuracy: 0.9474
 - Precision: 0.9459
 - Recall: 0.9722
 - F1 Score: 0.9589
-

2.2 Random Forest Classifier

- **Confusion Matrix:** $\begin{bmatrix} 39 & 3 \\ 3 & 69 \end{bmatrix}$
 - **Metrics:**
 - Accuracy: 0.9474
 - Precision: 0.9583
 - Recall: 0.9583
 - F1 Score: 0.9583
-

2.3 AdaBoost Classifier

- **Confusion Matrix:** $\begin{bmatrix} 38 & 4 \\ 2 & 70 \end{bmatrix}$
 - **Metrics:**
 - Accuracy: 0.9474
 - Precision: 0.9459
 - Recall: 0.9722
 - F1 Score: 0.9589
-

2.4 Voting Classifier

The Voting Classifier combines all three tuned models using **soft voting** and achieved the following:

- **Confusion Matrix:** $\begin{bmatrix} 38 & 4 \\ 2 & 70 \end{bmatrix}$
 - **Metrics:**
 - Accuracy: 0.9474
 - Precision: 0.9459
 - Recall: 0.9722
 - F1 Score: 0.9589
-

3. Declaring the Winner

Based on the evaluation metrics:

- **Decision Tree** and **AdaBoost** achieved the highest F1 Score (0.9589) and performed equally well.
- **Random Forest** achieved a slightly lower F1 Score (0.9583), though still highly competitive.

The **Decision Tree Classifier** is the **simplest** model with the best F1 score and is thus the most efficient in this scenario. However, the **Voting Classifier** is recommended for deployment because it combines multiple models, ensuring robustness and stability across predictions.

4. Conclusion

This process demonstrates the effectiveness of hyperparameter tuning using grid search and cross-validation. While the individual models performed well, the ensemble model (Voting Classifier) provides a balanced approach by leveraging the strengths of all models.